

class10

Michael Preston (PID: A53310268)

Class 10: Halloween Mini-Project

Exploratory Analysis of Halloween Candy

1. Importing candy data

```
#candy = read.csv("candy-data.csv", row.names=1)
library(readr)
# read in data and set column names
candy = read_csv("candy-data.csv", show_col_types = FALSE)
candy <- as.data.frame(candy)
rownames(candy) <- candy[[1]]
candy <- candy[, -1]
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0	0.732		0.860	66.97173		
3 Musketeers	0	1	0	0.604		0.511	67.60294		
One dime	0	0	0	0.011		0.116	32.26109		
One quarter	0	0	0	0.011		0.511	46.11650		
Air Heads	0	0	0	0.906		0.511	52.34146		
Almond Joy	0	1	0	0.465		0.767	50.34755		

```
# flextable allows for scrollable table in rendered doc
flextable::flextable(head(candy))
```

chocolate	fruity	caramel	peanut	almond	monkey	nougat	crisped	rice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	1	0	0	1	1	0	
1	0	0	0	0	1	0	0	0	0	1	1	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	
0	1	0	0	0	0	0	0	0	0	0	0	0	
1	0	0	1	0	0	0	0	0	0	0	1	0	

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Alternative way to run functions with ‘pipe’ in dplyr. This can be more readable for multi-step operations

```
candy |> select(winpercent) |> sum()
```

[1] 4276.925

Q1. How many different candy types are in this dataset?

There are 85 different candy types

Q2. How many fruity candy types are in the dataset?

There are 38 fruity candy types

2. What is your favorite candy?

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

Reese's Peanut Butter cup: 84

Q4. What is the winpercent value for "Kit Kat"?

77

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

50

try out skimr package

```
library("skimr")
skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

winpercent

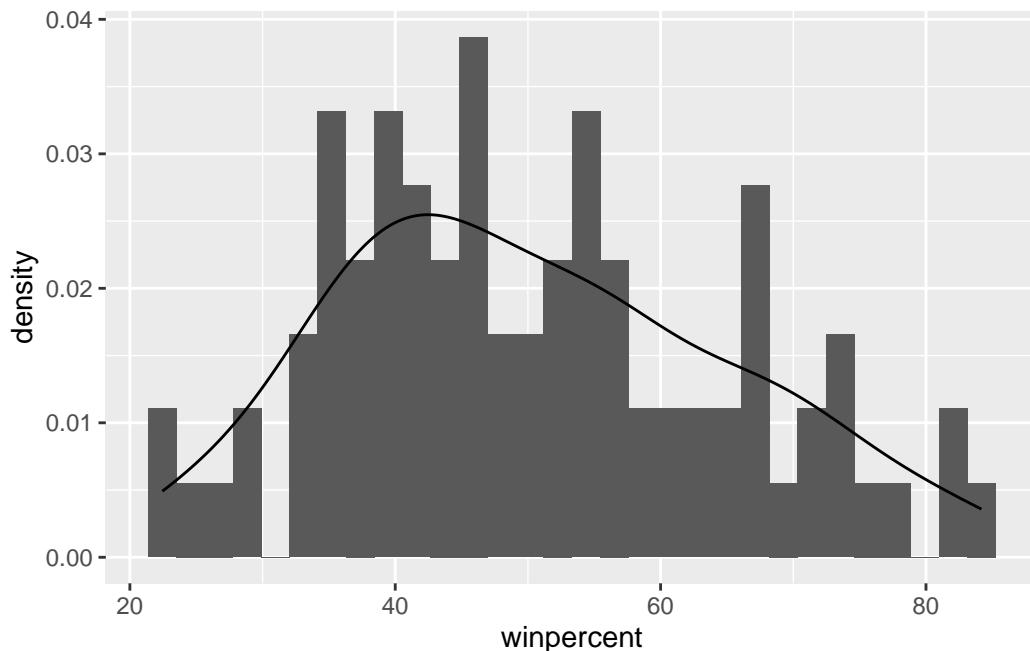
Q7. What do you think a zero and one represent for the candy\$chocolate column?

Logical/boolean indicator of whether the candy contains chocolate.

Q8. Plot a histogram of winpercent values

```
library("ggplot2")
ggplot(candy, aes(x = winpercent)) +
  geom_histogram(aes(y = after_stat(density))) +
  geom_density()
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.



Q9. Is the distribution of winpercent values symmetrical?

Somewhat right skewed

Q10. Is the center of the distribution above or below 50%?

The median is below, while the mean is slightly above.

mean: 50.3167638 median: 47.829754

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Higher

```
mean(candy[candy$fruity==1,]$winpercent)
```

[1] 44.11974

```
mean(candy[candy$chocolate==1,]$winpercent)
```

[1] 60.92153

Q12. Is this difference statistically significant?

Yes

```
t.test(candy[candy$fruity==1,]$winpercent, candy[candy$chocolate==1,]$winpercent)
```

Welch Two Sample t-test

```
data: candy[candy$fruity == 1, ]$winpercent and candy[candy$chocolate == 1, ]$winpercent
t = -6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-22.15795 -11.44563
sample estimates:
mean of x mean of y
44.11974 60.92153
```

3. Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters

```
candy |> arrange(winpercent) |> head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0	1		0.197
Boston Baked Beans	0	0	0	1		0.313
Chiclets	0	0	0	1		0.046
Super Bubble	0	0	0	0		0.162
Jawbusters	0	1	0	1		0.093
	percent					pricepercent
Nik L Nip	22.44534					0.976
Boston Baked Beans	23.41782					0.511
Chiclets	24.52499					0.325
Super Bubble	27.30386					0.116
Jawbusters	28.12744					0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

Q14. What are the top 5 all time favorite candy types out of this set?

Reese's Peanut Butter cup, Reese's Miniatures, Twix, Kit Kat, and Snickers.

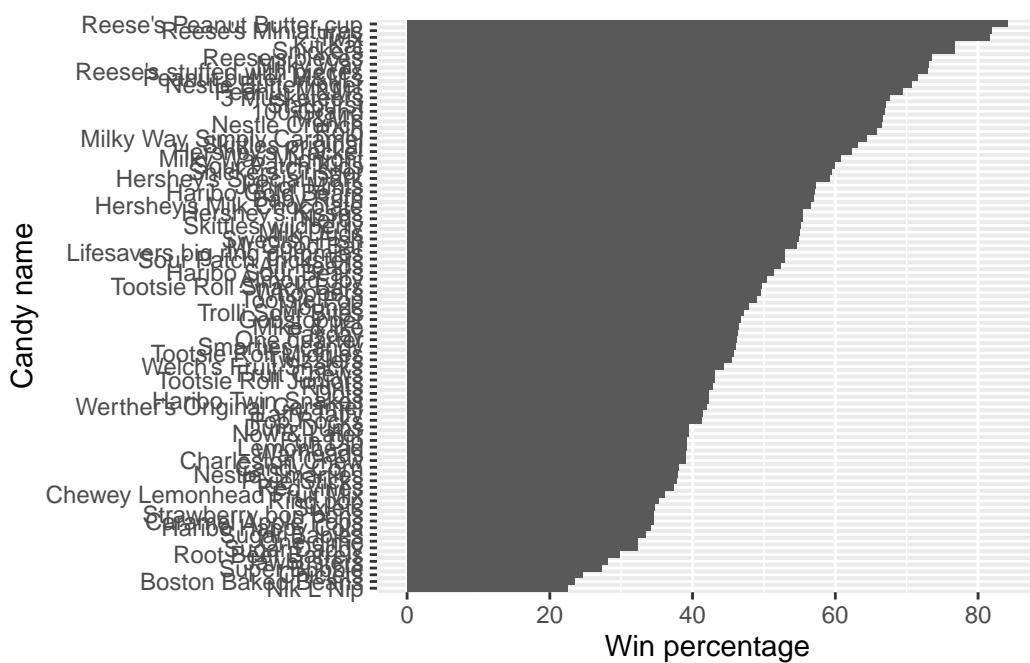
```
candy |> arrange(winpercent) |> tail(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crisped	rice	wafer	hard	bar	pluribus
Snickers	0	0	1		0	0.546
Kit Kat	1	0	1		0	0.313

	1	0	1	0	0.546
Reese's Miniatures	0	0	0	0	0.034
Reese's Peanut Butter cup	0	0	0	0	0.720
Snickers	0.651	76.67378			
Kit Kat	0.511	76.76860			
Twix	0.906	81.64291			
Reese's Miniatures	0.279	81.86626			
Reese's Peanut Butter cup	0.651	84.18029			

Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col() +
  ylab('Candy name') +
  xlab('Win percentage')
```



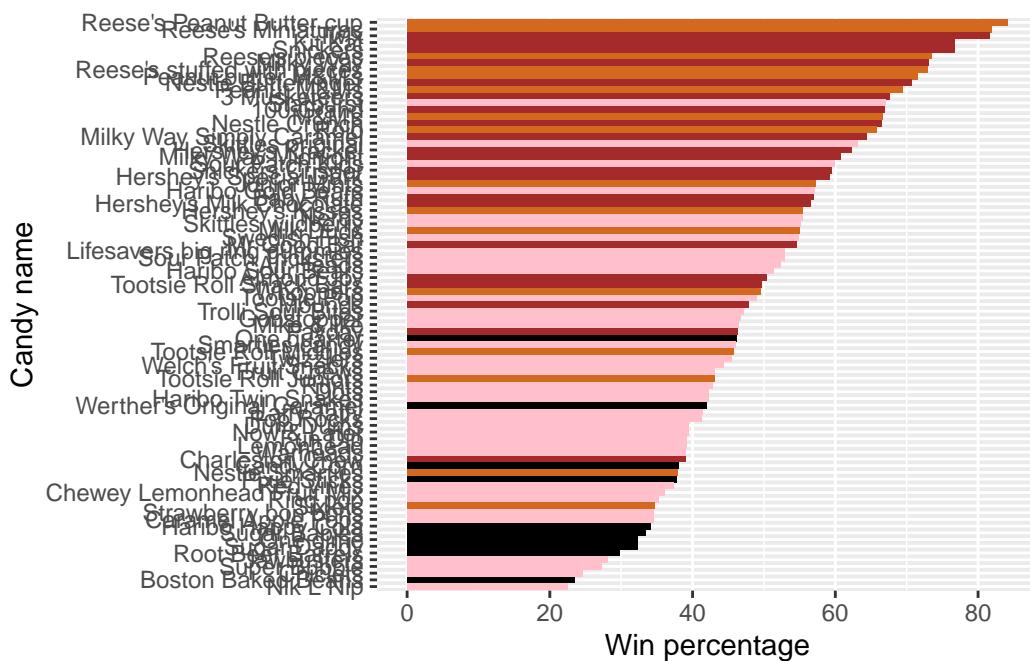
Time to add some useful color

```

my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols) +
  ylab('Candy name') +
  xlab('Win percentage')

```



Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?

Starburst

```

temp <- arrange(candy[as.logical(candy$fruity),], winpercent)
rownames(temp)[nrow(temp)]

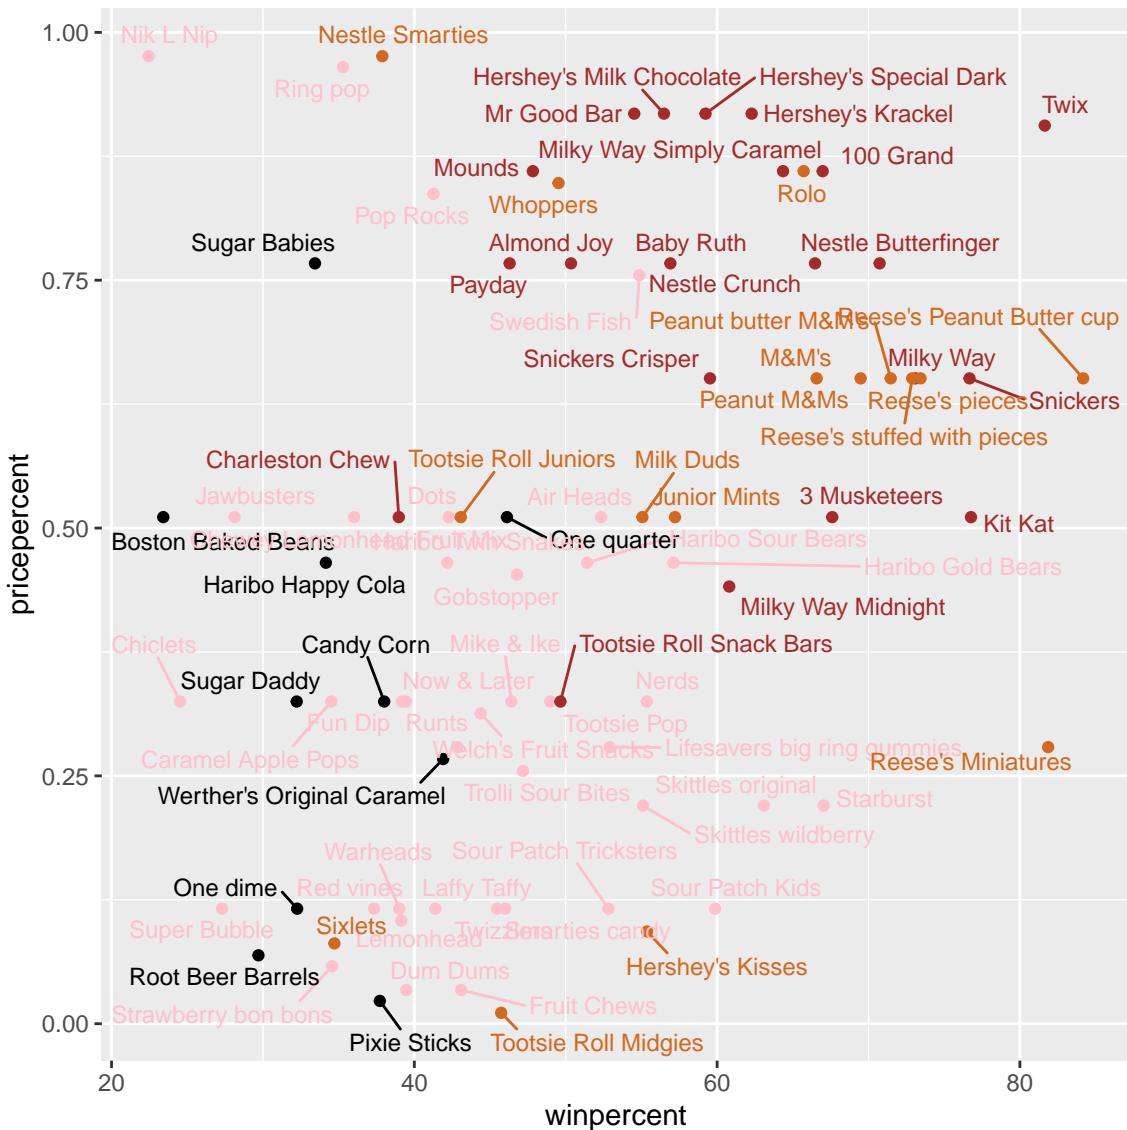
```

[1] "Starburst"

4. Taking a look at pricepercent

```
library(ggrepel)

# How about a plot of price vs win
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 20)
```



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reeses Miniatures is a good balance

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

5 most expensive candy types: “Nik L Nip”, “Nestle Smarties”, “Ring pop”, “Hershey’s Krackel”, “Hershey’s Milk Chocolate”

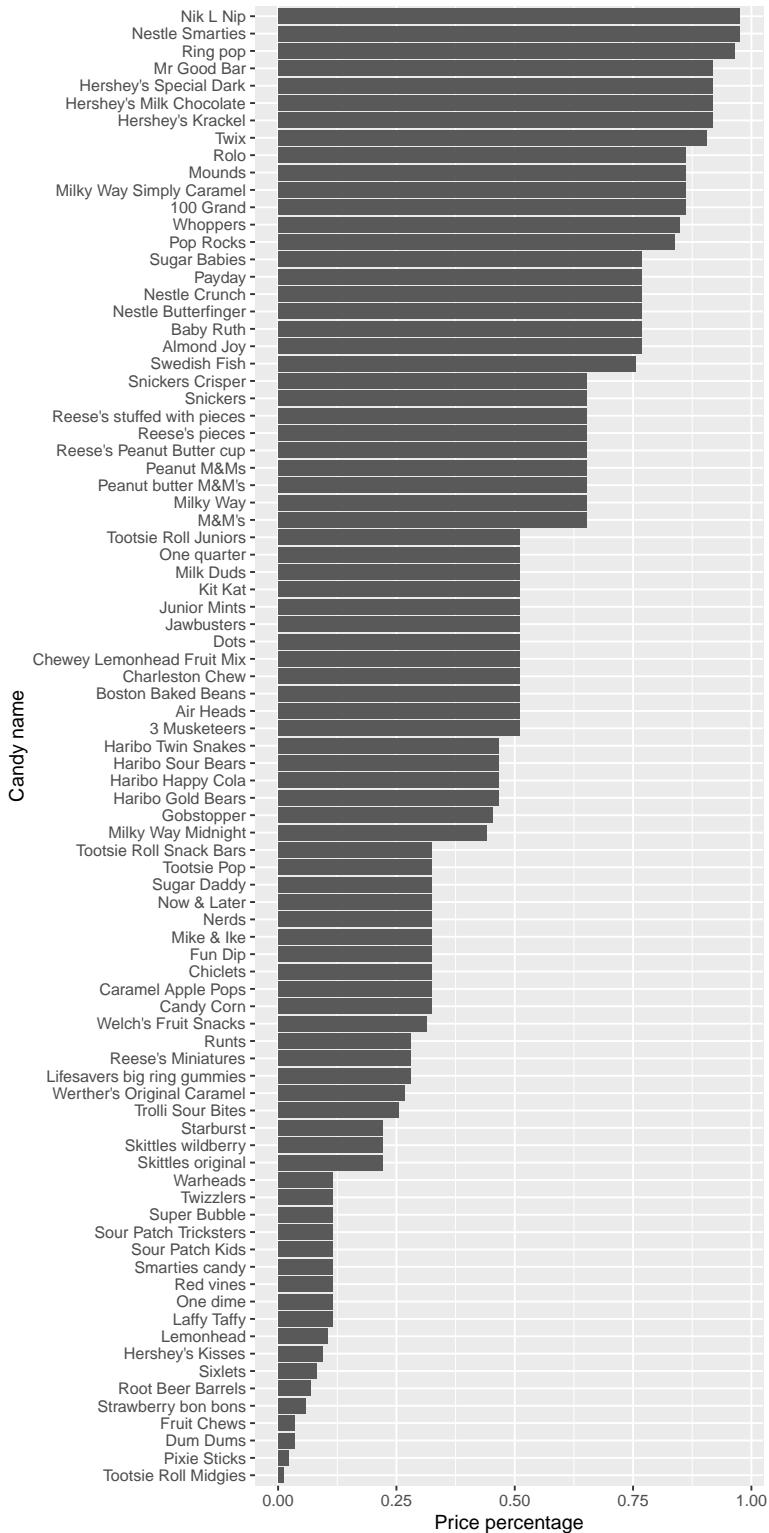
Least popular: Nik L Nip

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,], n=5 )
```

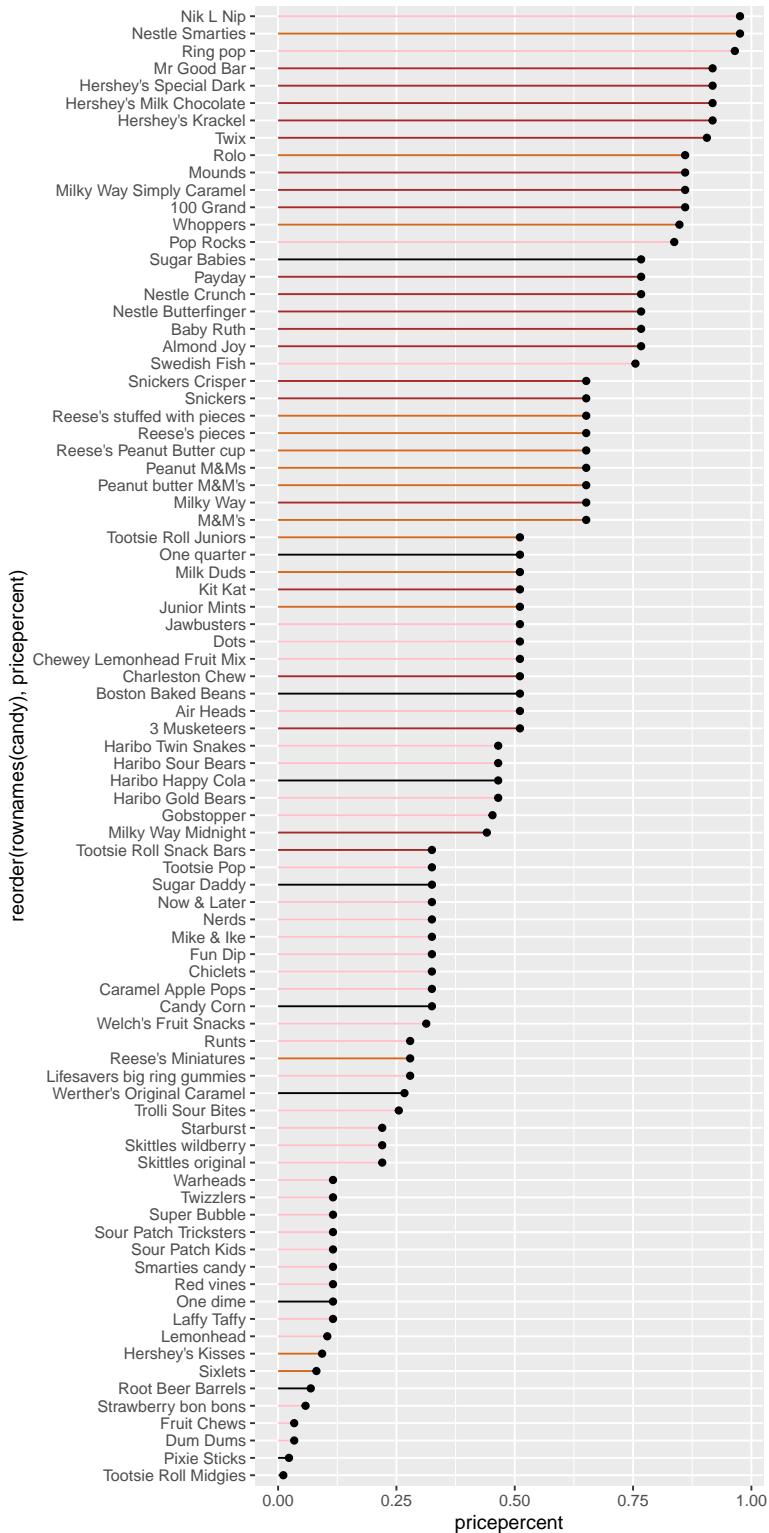
	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0	0	0	0
Nestle Smarties	1	0	0	0	0	0
Ring pop	0	1	0	0	0	0
Hershey's Krackel	1	0	0	0	0	0
Hershey's Milk Chocolate	1	0	0	0	0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0	1		0.197
Nestle Smarties	0	0	0	1		0.267
Ring pop	0	1	0	0		0.732
Hershey's Krackel	1	0	1	0		0.430
Hershey's Milk Chocolate	0	0	1	0		0.430
	pricepercent	winpercent				
Nik L Nip	0.976	22.44534				
Nestle Smarties	0.976	37.88719				
Ring pop	0.965	35.29076				
Hershey's Krackel	0.918	62.28448				
Hershey's Milk Chocolate	0.918	56.49050				

Q21. Make a barplot again with geom_col() this time using pricepercent and then improve this step by step, first ordering the x-axis by value and finally making a so called “dot chat” or “lollipop” chart by swapping geom_col() for geom_point() + geom_segment()

```
ggplot(candy) +  
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +  
  geom_col() +  
  ylab('Candy name') +  
  xlab('Price percentage')
```



```
# Make a lollipop chart of pricepercent
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy), pricepercent),
                    xend = 0), col=my_cols) +
  geom_point()
```

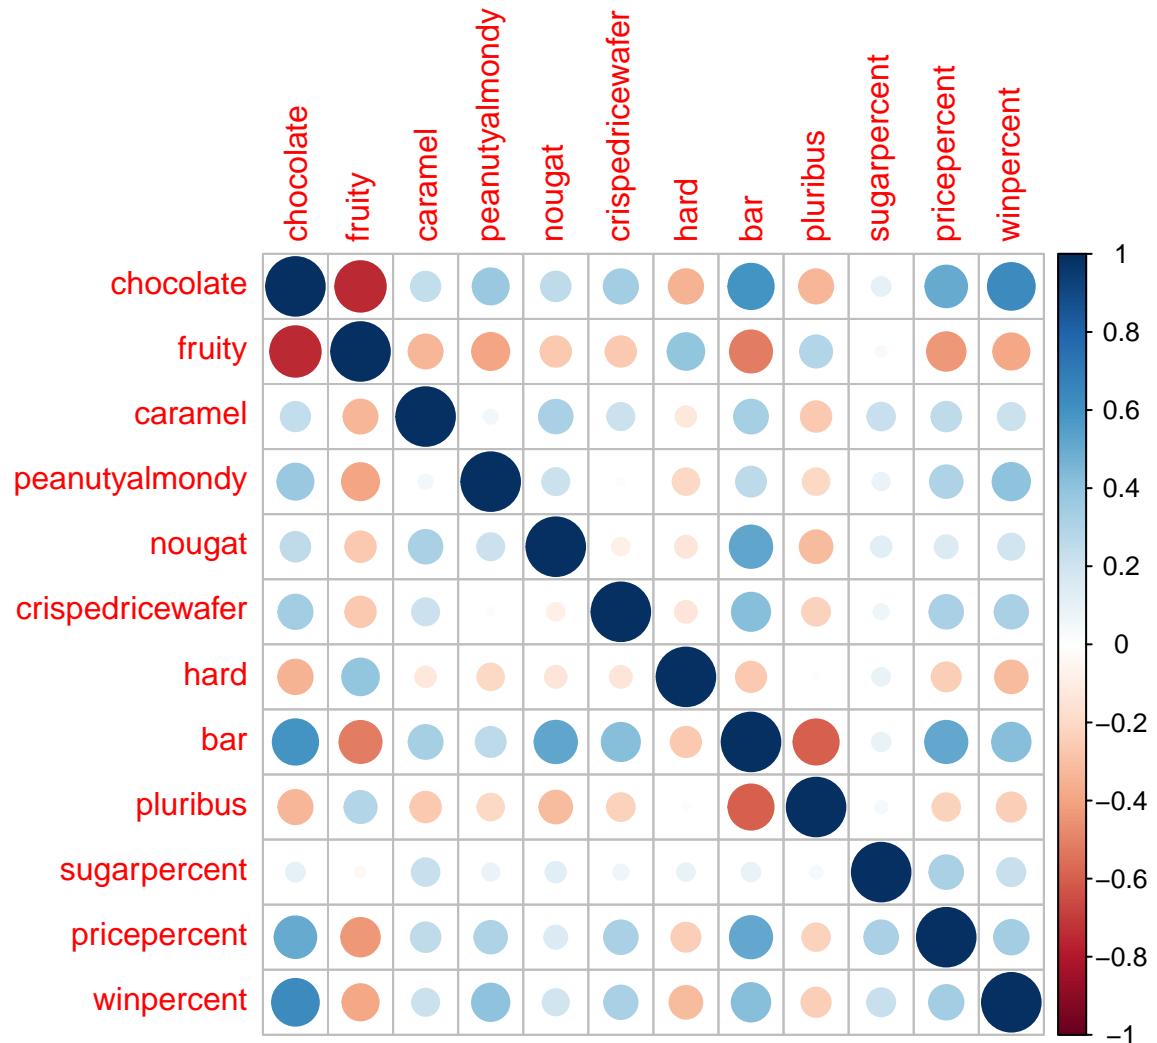


5 Exploring the correlation structure

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and Fruity have the greatest negative correlation.

Q23. Similarly, what two variables are most positively correlated?

winpercent and chocolate

6. Principal Component Analysis

```
# Define plotting functions for PCA results

plot_pca_scatter12 <- function(pca){
  # Plot "PC Plot" / "score plot" i.e. scatterplot of PC1 v. PC2
  ggplot(pca$x) +
    aes(x = PC1, y = PC2) +
    geom_point()
}

plot_pca_scree <- function(pca){
  # Create scree plot

  # Compute variance explained
  v <- round(pca$sdev^2/sum(pca$sdev^2) * 100)
  variance_df <- data.frame(
    PC = factor(paste0("PC", 1:length(v)), levels = paste0("PC", 1:length(v))),
    Variance = v
  )

  # Plot
  ggplot(variance_df) +
    aes(x = PC, y = Variance) +
    geom_col(fill = "steelblue") +
    xlab("Principal component") +
    ylab("Percent variation") +
    theme_bw() +
    theme(axis.text.x = element_text(angle = 0))
}

plot_pca_loading <- function(pca){
  # plot PCA loadings
```

```

ggplot(pca$rotation) +
  aes(x = PC1,
  y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
}

```

```

plot_pca_results <- function(pca){
  print(plot_pca_scatter12(pca))
  print(plot_pca_scree(pca))
  print(plot_pca_loading(pca))
}

```

```

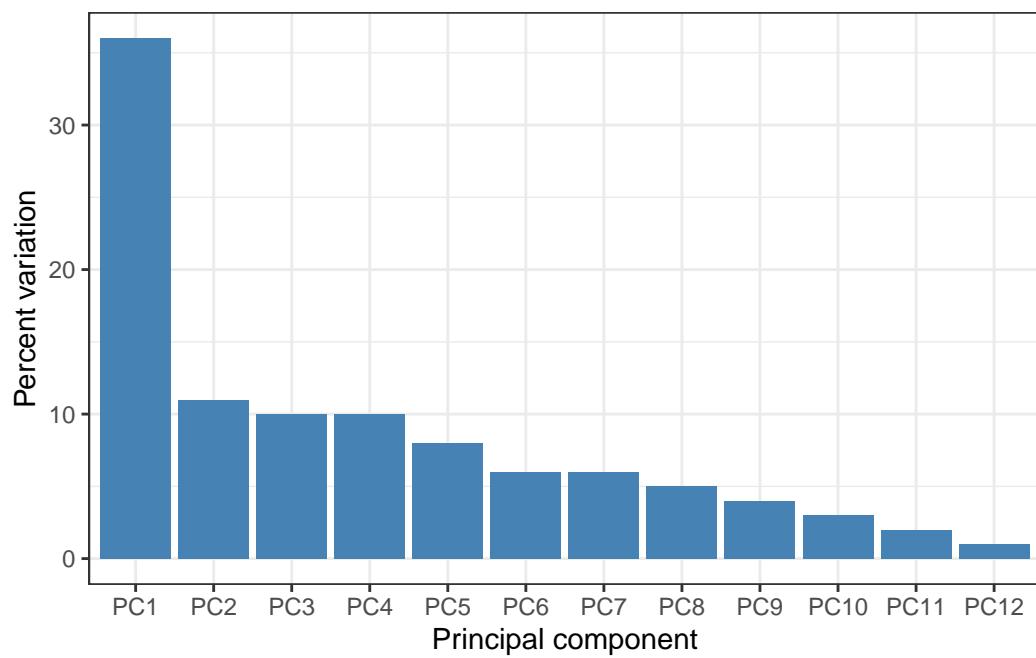
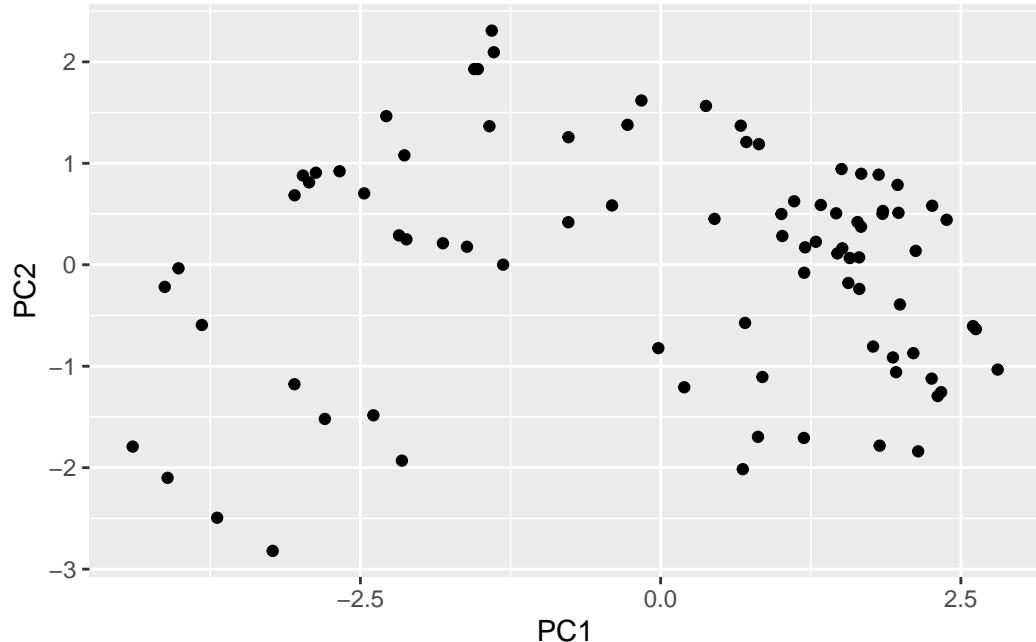
# Perform PCA and show results
candy.pca <- prcomp(candy, scale=T)
summary(candy.pca)

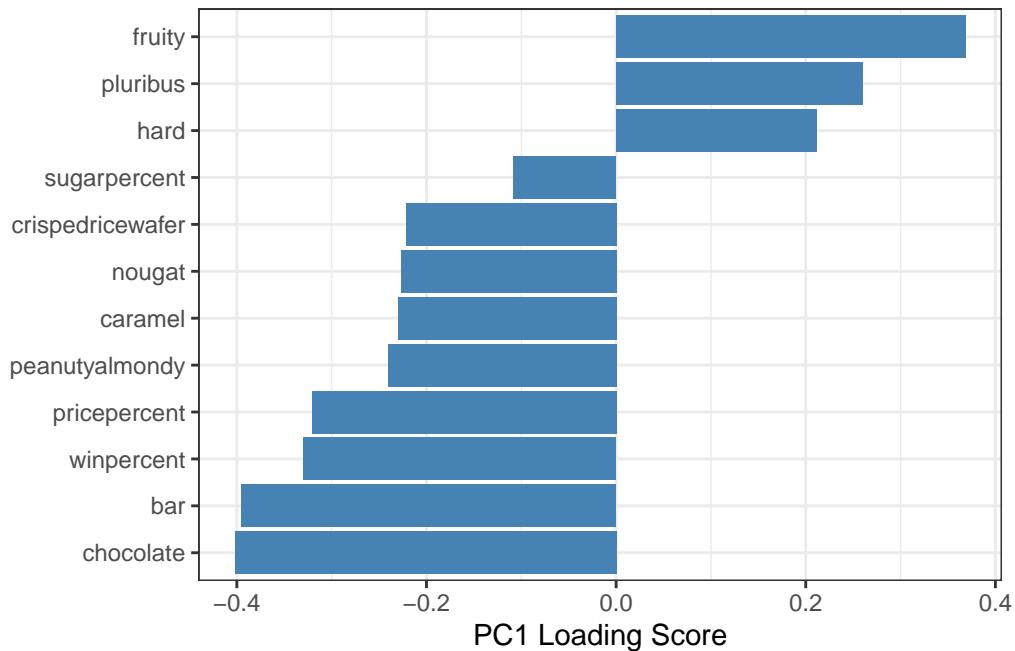
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

```
plot_pca_results(candy.pca)
```



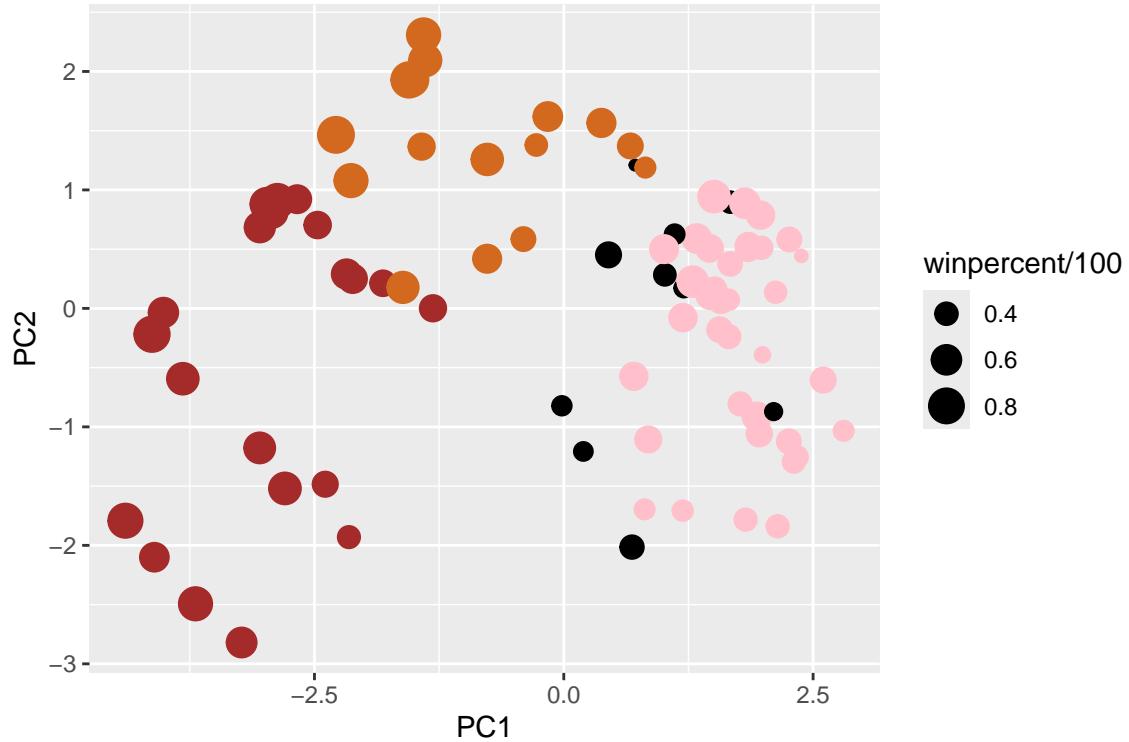


```
# repeat, plotting size based on winpercent

# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, candy.pca$x[,1:3])

# plot
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p
```

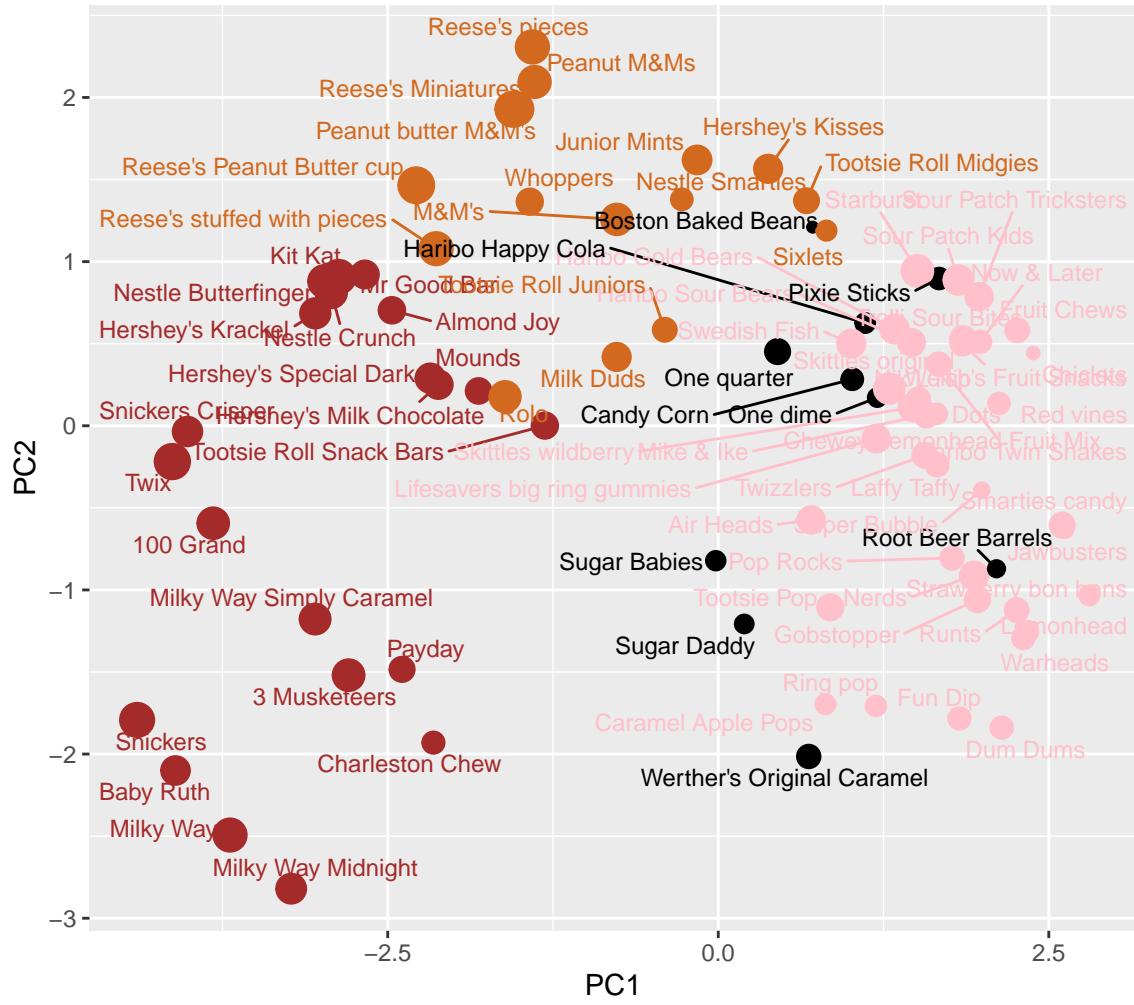


Use ggrepel package and the function `ggrepel::geom_text_repel()` to label up the plot with non overlapping candy names

```
p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 30) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
       caption="Data from 538")
```

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown), fruity (pink)



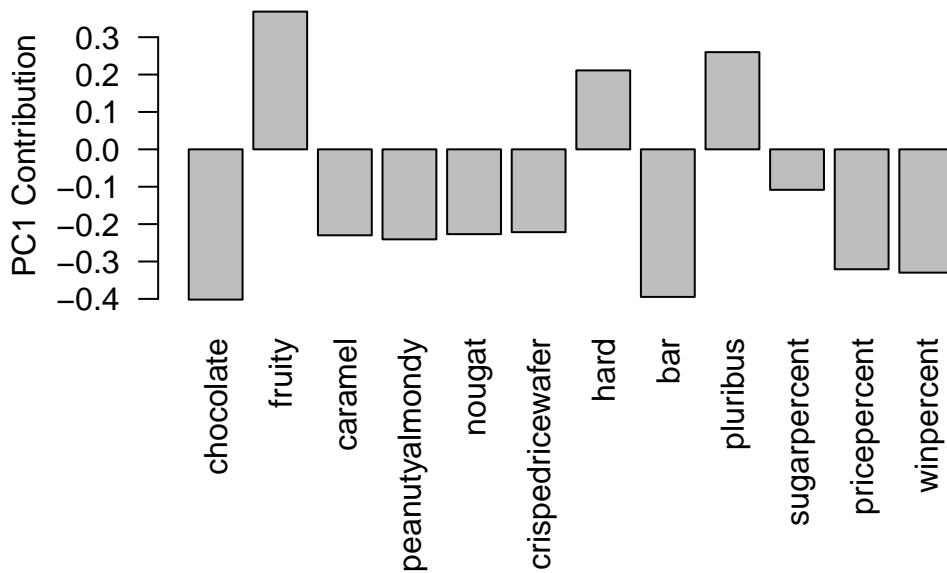
Data from 538

One more plotting option with gplotly. This is an interactive plot, so you can get info on each data point by hovering the mouse over it

```
# library(plotly)
# ggplotly(p)
```

Check out loadings

```
par(mar=c(8,4,2,2))
barplot(candy.pca$rotation[,1], las=2, ylab="PC1 Contribution")
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard, and pluribus. Yes, these features describe the candies with the highest values along PC1 (e.g., skittles, warheads, lemonheads)